

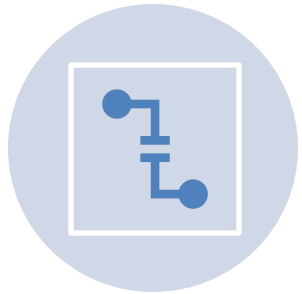
# Neural Networks

## *Multi-Layer Networks*

CSCI 4850/5850



# Linear = Limited



Single-layer networks are inherently constrained in the kinds of input-output mappings (i.e. functions) that they can produce: **regardless of the learning algorithm used!**



In particular, if input vectors are not **linearly independent**, then the range of allowable outputs is limited.



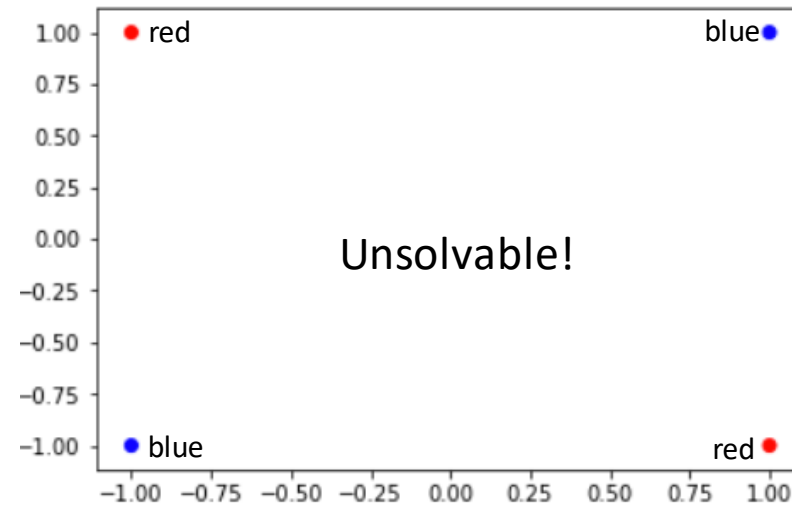
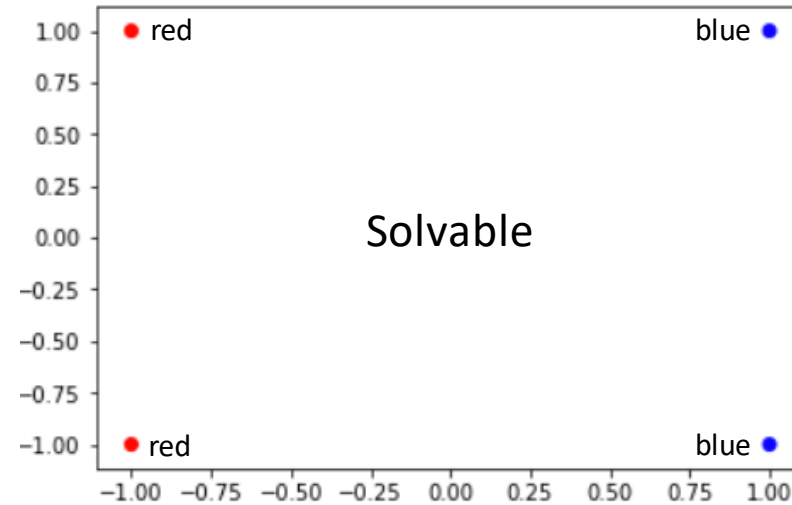
For perfect performance, the target,  $t$ , for input vector,  $p$ , must be...



For classification, this is the **linear separability** constraint

$$t^p = g(\mathbf{w} \cdot \mathbf{x}^p) = g\left(\sum_{n=1}^N v_n(\mathbf{w} \cdot \mathbf{x}^n)\right) = g\left(\sum_{n=1}^N v_n g^{-1}(t^n)\right) \text{ where } \mathbf{x}^p = \sum_{n=1}^N v_n \mathbf{x}^n$$

# Constraints of Single-layer Nets



Classes:

Blue = 0

Red = 1

Data (Top) :

$[-1 \ -1] = 1$

$[-1 \ 1] = 1$

$[1 \ -1] = 0$

$[1 \ 1] = 0$

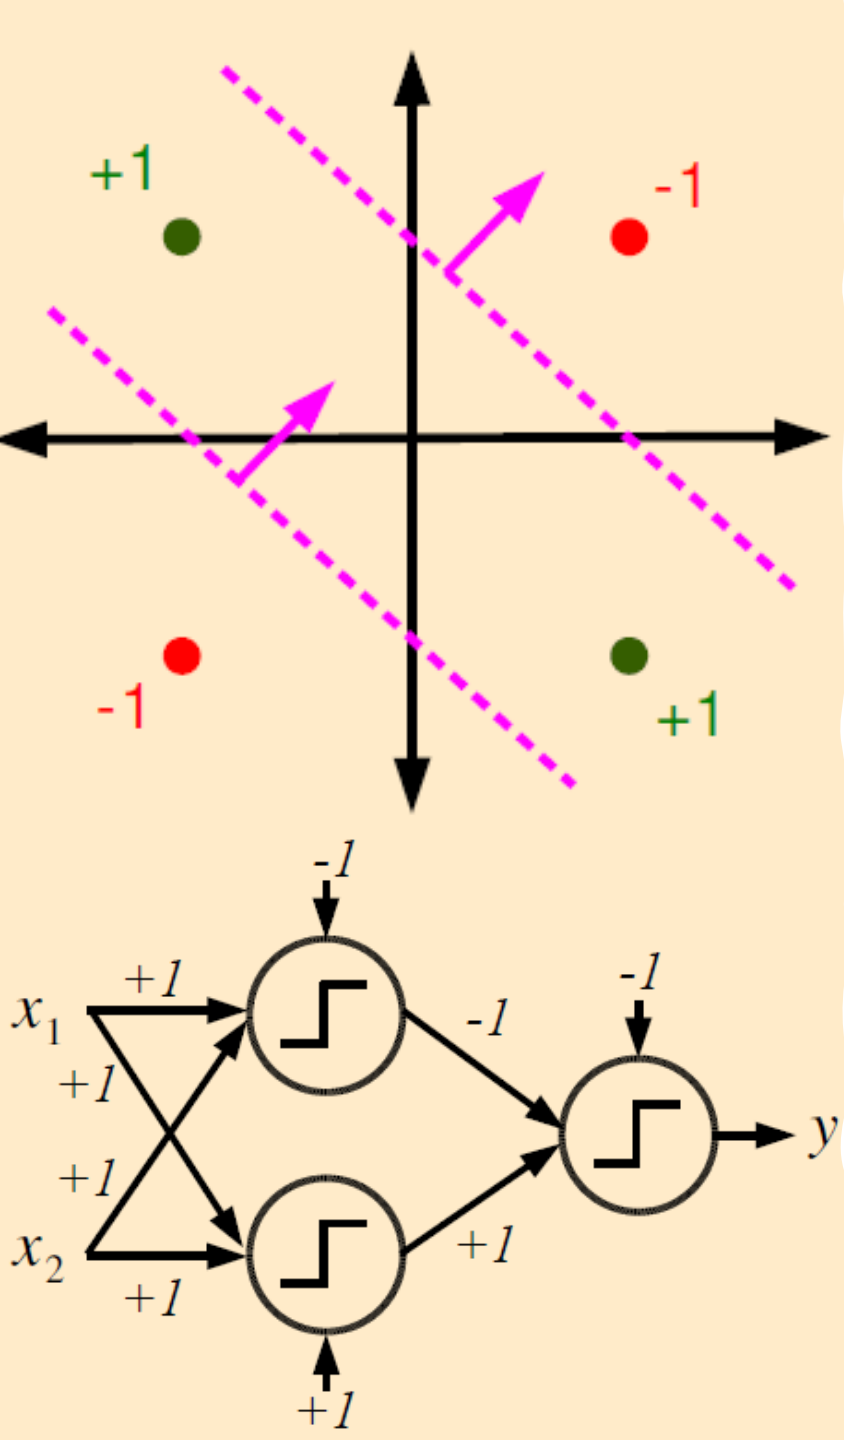
Data (Bottom) :

$[-1 \ -1] = 0$

$[-1 \ 1] = 1$

$[1 \ -1] = 1$

$[1 \ 1] = 0$

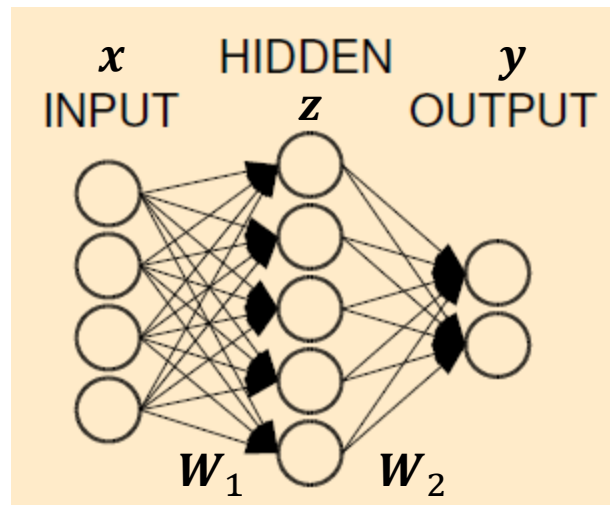


# The Utility of Multiple Layers

- Multilayer networks can solve **non-linearly separable** problems by combining multiple partitions of the input vector space
- Different partitions are computed using different **hidden units**

# Key: Linear/Non-linear Activation Functions

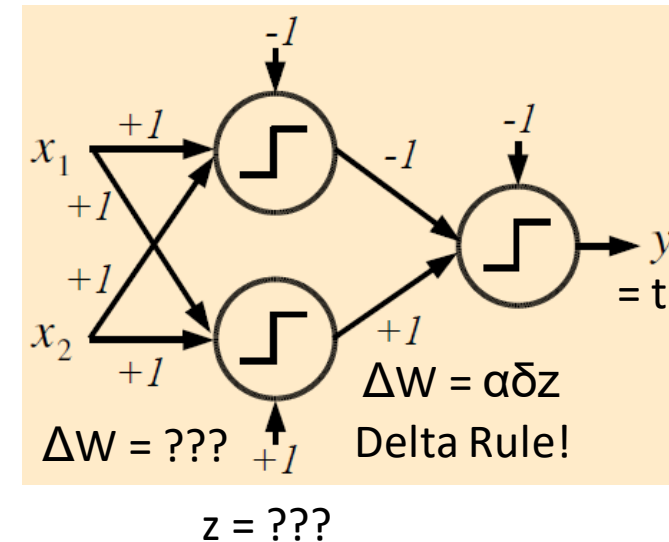
- Even a multilayer network will provide **no benefit** if the hidden units all utilize a **linear/affine activation function**



$$y = W_2 z = W_2(W_1 x) = (W_2 W_1)x = W_m x$$

# What to do?

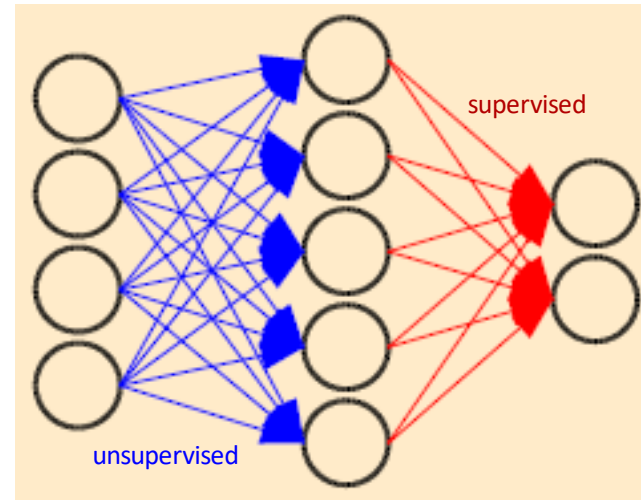
- For single-layer networks, the **target activation** for **output units** is known from the *training data*
- But, it's not clear what the activations of the hidden units need to be...



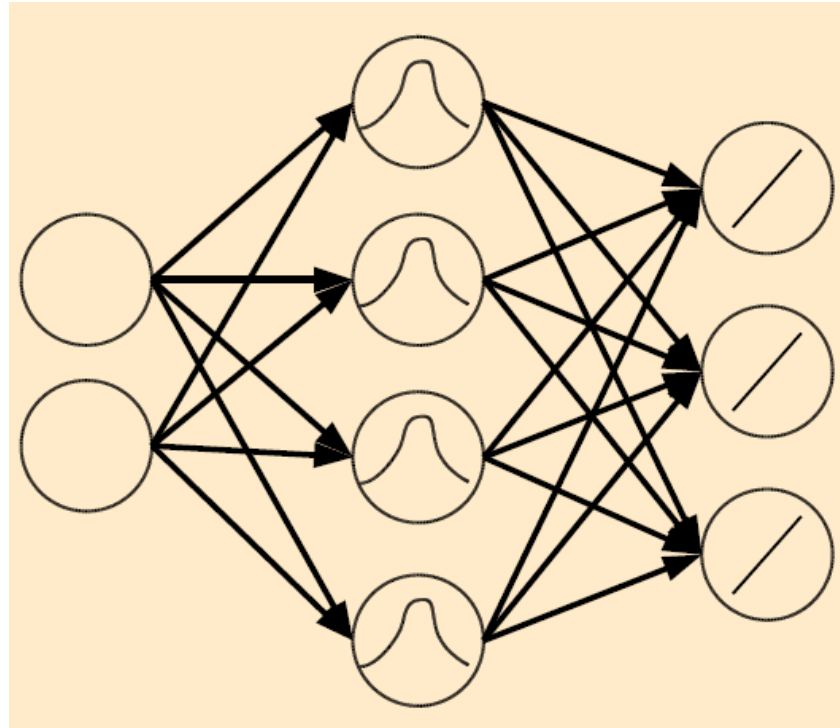
# One Potential Solution

Apply *unsupervised* learning

Different kind of learning rule that tries to adapt the weights into something useful without *target* information



# Radial Basis Function Networks



Counterpropagation Networks - Hecht-Neilsen, 1987

$$y_k(\mathbf{x}) = \sum_{j=1}^M w_{kj} \phi_j(\mathbf{x})$$

$$\phi_j(\mathbf{x}) = \exp\left(-\frac{\|\mathbf{x} - \boldsymbol{\mu}_j\|^2}{2\sigma_j^2}\right)$$

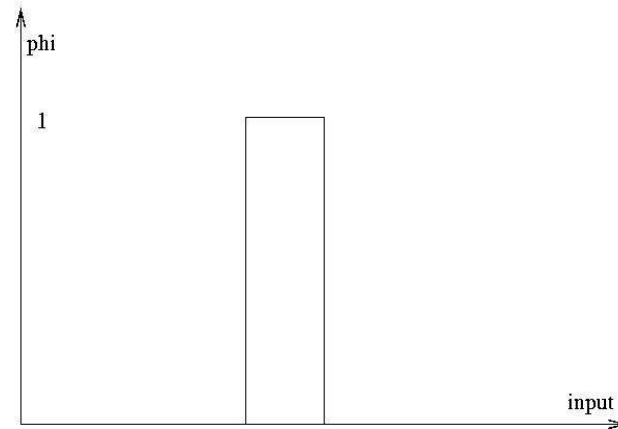
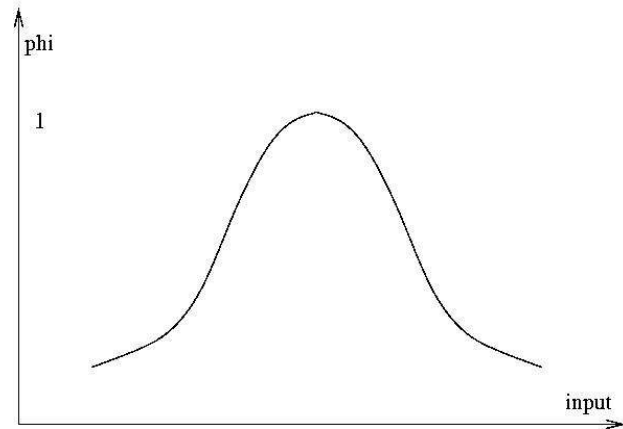
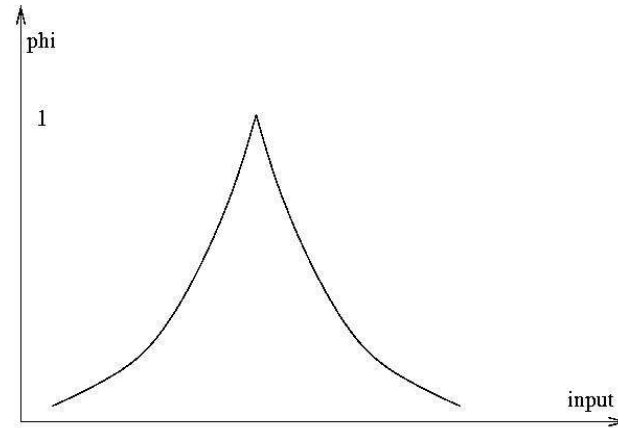
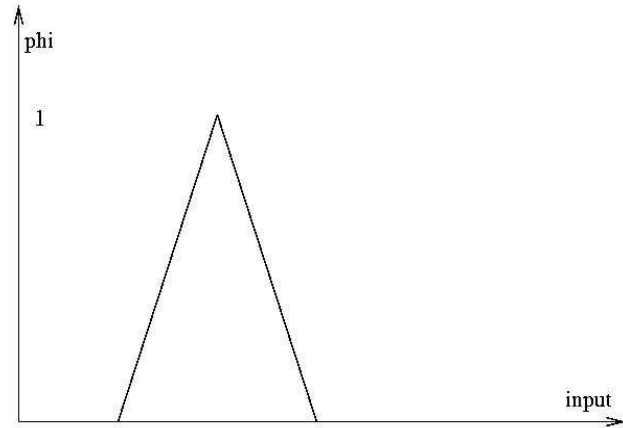
Can be augmented with neighborhood topologies as well!

Mixes unsupervised with supervised approaches

There are learning rules for the other parameters (sigma)...



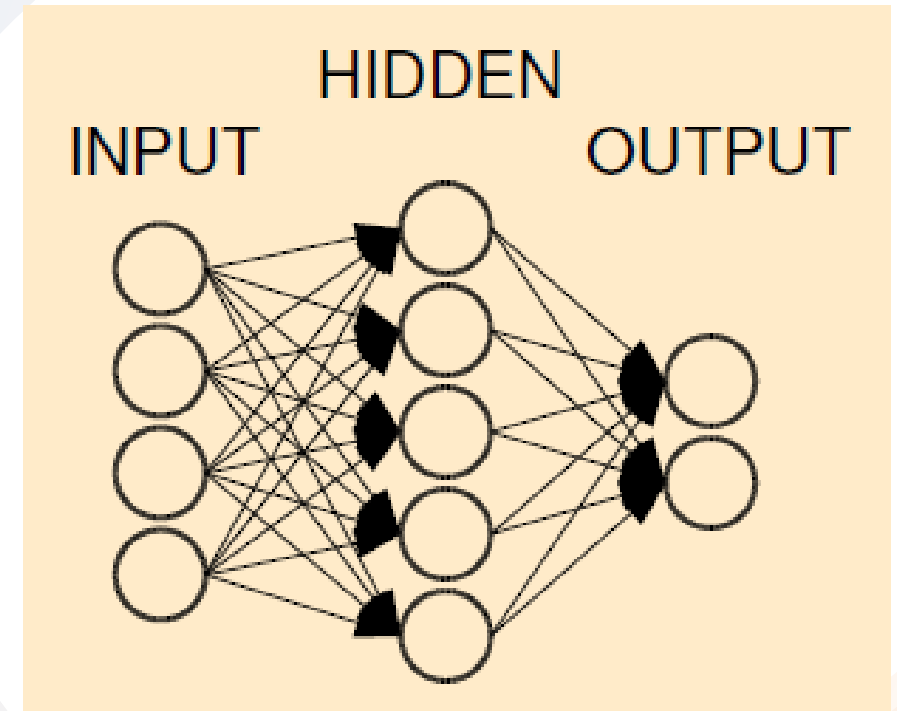
# Interpolation – Radial Basis Functions



$$\phi_j(\mathbf{x}) = \phi(\|\mathbf{x} - \boldsymbol{\mu}_j\|)$$

# Universal Approximation

- It has been proven that given a **two-layer network** and **enough *nonlinear* hidden units**, essentially *any* function can be approximated to an arbitrary degree of precision
- Hornik, Stinchcombe, and White, 1989
- A nonlinear activation function is required...



# Can we obtain information about the error to help us?

- As we saw before with the SSE of a single-layer network, the training data (input and output vectors) are *fixed*.
- Hidden unit activations are dependent on the data in the same way as output units.
- Thus, the weights are the only real parameters that have an influence on the activations!

$$y = F(x) \qquad \frac{\partial y}{\partial w_{ij}} = \frac{\partial F(x)}{\partial w_{ij}}$$

# It's "complex" but not incomputable...

The chain rule makes it all possible...

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial}{\partial w_{ij}} \frac{1}{2} \sum_{k=1}^c (y_k - t_k)^2 = \sum_{k=1}^c (y_k - t_k) \frac{\partial y_k}{\partial w_{ij}}$$

We'll derive it later, for now let's just see the result...

$$\Delta w_{ij} \propto -\frac{\partial E}{\partial w_{ij}}$$

For output units...

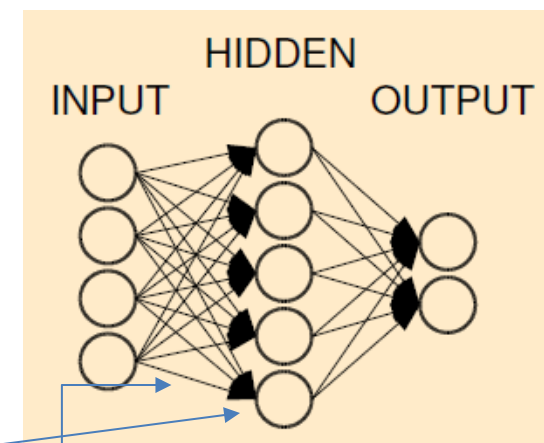
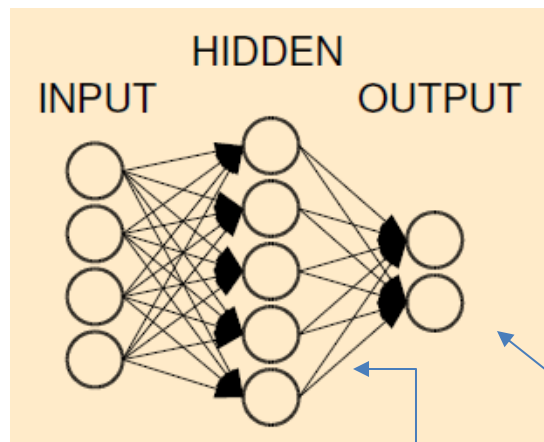
$$\delta_j = g'(net_j)(a_j - t_j)$$

For hidden units...

$$\delta_j = g'(net_j) \sum_k w_{jk} \delta_k$$

$$\Delta w_{ij} \propto -\delta_j a_i$$

Note that these equations apply for any weight connecting unit  $i$  to unit  $j$ .



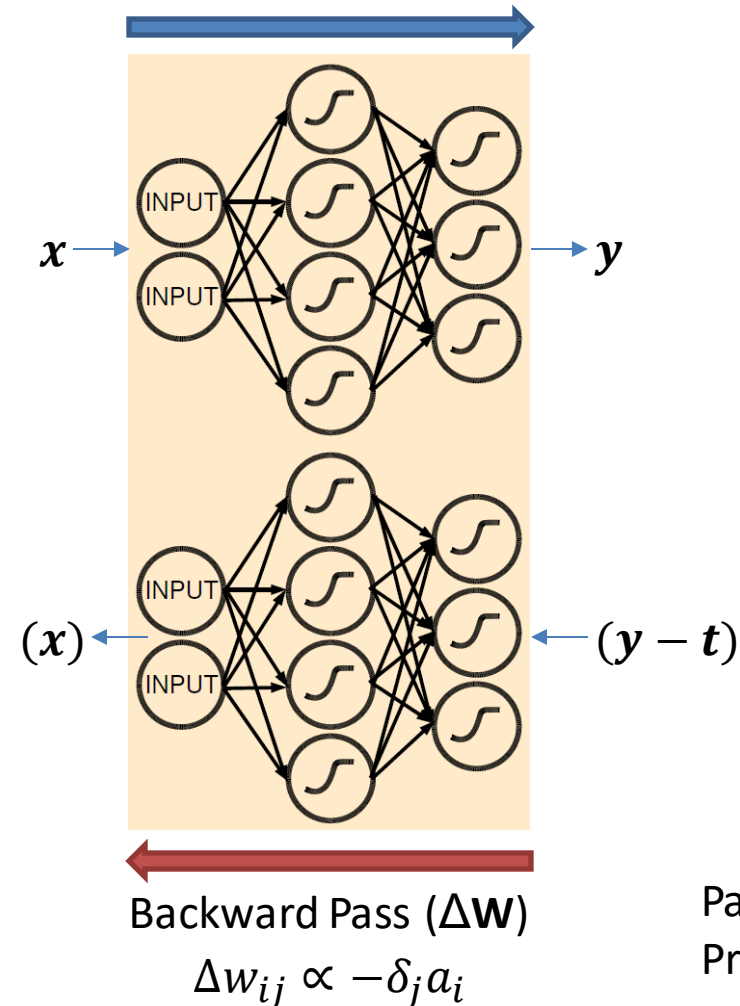
# Generalized Delta Rule

- Also known as **error backpropagation** or “**backprop**”
- Two step process:
  - **Prediction:** the inputs are provided and information flows to calculate the output activations (forward pass)
  - **Fitting:** errors are calculated at the output layer and information flows in reverse to calculate the weight updates (backward pass)
- Issue: biological plausibility
  - Real neurons do not pass information backward across synapses...

Output unit:  $\delta_j = g'(net_j)(a_j - t_j)$

Hidden unit:  $\delta_j = g'(net_j) \sum_k w_{jk} \delta_k$

$$a_i = g(net_i) \quad net_i = w_{i0} + \sum_j w_{ij} a_j$$



Parallel Distributed Processing (Rumelhart and McClelland, 1986)

# Backprop Derivation

- The Chain Rule is the Key

$$\frac{df}{dx} = \frac{df}{du} \frac{du}{dx}$$

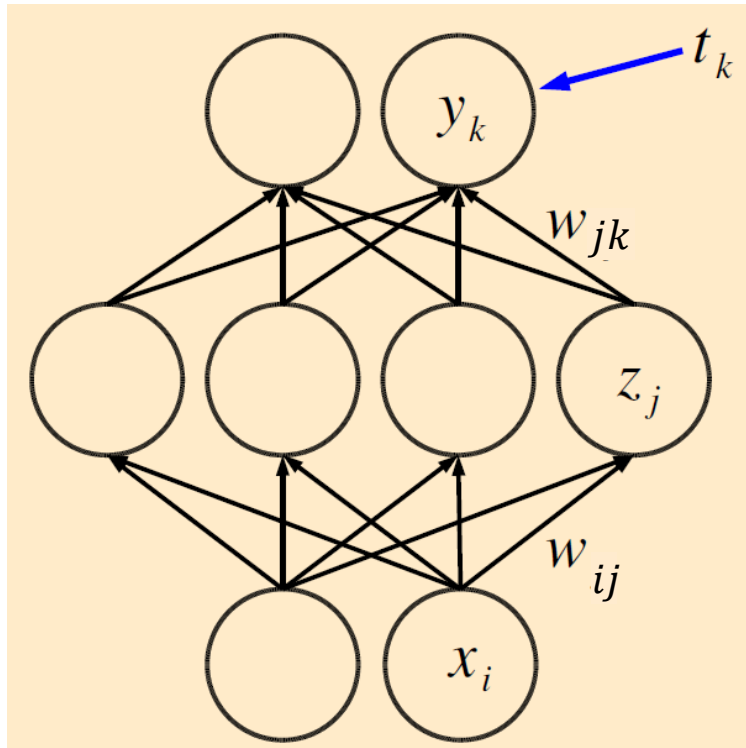
$$\frac{\partial f}{\partial x} = \sum_{i=1}^n \frac{\partial f}{\partial g_i} \frac{\partial g_i}{\partial x}$$

$$E = \sum_{n=1}^N E^n$$

$$E^n = \frac{1}{2} \sum_{k=1}^c (y_k^n - t_k^n)^2$$

$$\delta_j \equiv \frac{\partial E^n}{\partial net_j}$$

$$f(x) = f(g_1(x), g_2(x), \dots, g_n(x))$$



$C$  output units

$M$  hidden units

$D$  input units

# Delta for an Output Unit

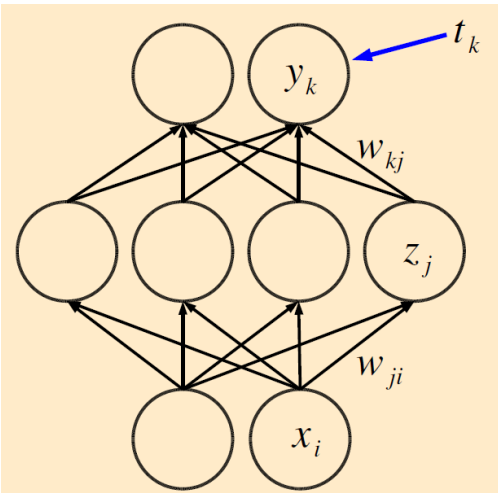
$$\delta_k = \frac{\partial E^n}{\partial net_k} = \frac{\partial}{\partial net_k} \left( \frac{1}{2} \sum_{c=1}^C (y_c^n - t_c^n)^2 \right)$$

$$= \frac{1}{2} \frac{\partial}{\partial net_k} (y_k^n - t_k^n)^2$$

$$= (y_k^n - t_k^n) \frac{\partial}{\partial net_k} (y_k^n - t_k^n)$$

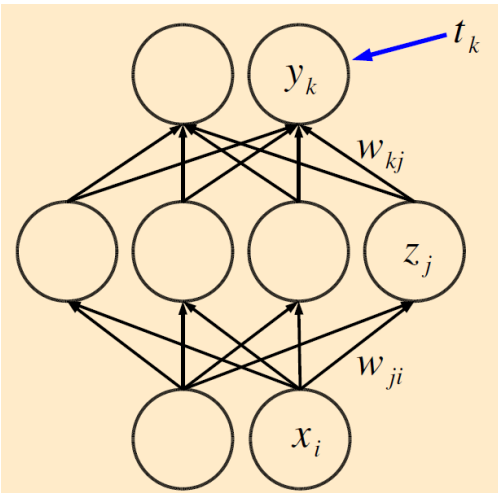
$$= (y_k^n - t_k^n) \frac{\partial y_k^n}{\partial net_k}$$

$$= (y_k^n - t_k^n) \frac{\partial g(net_k)}{\partial net_k} = (y_k^n - t_k^n) g'(net_k)$$



# Rederived Delta Rule

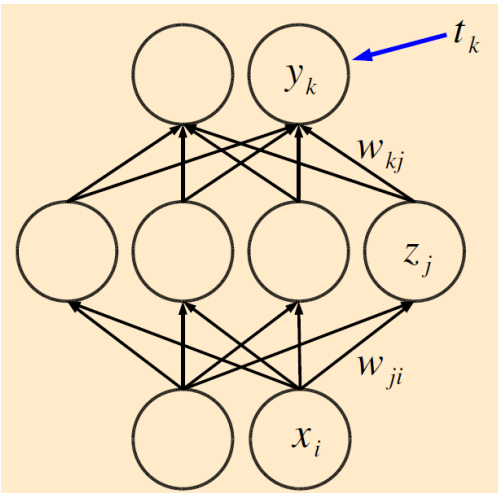
$$\begin{aligned}-\frac{\partial E^n}{\partial w_{jk}} &= -\frac{\partial E^n}{\partial net_k} \frac{\partial net_k}{\partial w_{jk}} \\&= -\delta_k \frac{\partial net_k}{\partial w_{jk}} \\&= -\delta_k \frac{\partial}{\partial w_{jk}} \left( \sum_{m=1}^M w_{mk} z_m^n \right) \\&= -\delta_k z_j^n \\&= -(y_k^n - t_k^n) g'(net_k) z_j^n \\&= -g'(net_k) (y_k^n - t_k^n) z_j^n\end{aligned}$$





# Gradient for a Hidden Unit

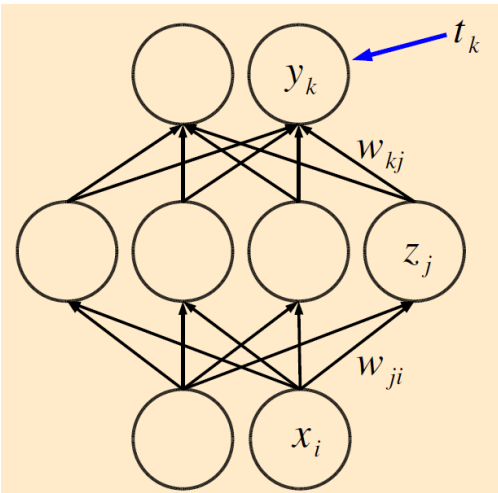
$$\begin{aligned} -\frac{\partial E^n}{\partial w_{ij}} &= -\frac{\partial E^n}{\partial net_j} \frac{\partial net_j}{\partial w_{ij}} \\ &= -\delta_j \frac{\partial}{\partial w_{ij}} \left( \sum_{d=1}^D w_{dj} x_d^n \right) \\ &= -\delta_j x_d^n \end{aligned}$$



OK, but what is the delta here?

# Delta for a Hidden Unit

$$\begin{aligned}
 \delta_j &= \frac{\partial E^n}{\partial net_j} = \frac{\partial E^n}{\partial z_j^n} \frac{\partial z_j^n}{\partial net_j} = \frac{\partial E^n}{\partial z_j^n} g'(net_j) \\
 &= g'(net_j) \sum_{c=1}^C \frac{\partial E^n}{\partial net_c} \frac{\partial net_c}{\partial z_j^n} \\
 &= g'(net_j) \sum_{c=1}^C \delta_c \frac{\partial}{\partial z_j^n} \left( \sum_{m=1}^M w_{mc} z_m^n \right) \\
 &= g'(net_j) \sum_{c=1}^C \delta_c w_{jc}
 \end{aligned}$$



So, now we can complete the gradient from before...

$$-\frac{\partial E^n}{\partial w_{ij}} = -\delta_j x_i^n$$

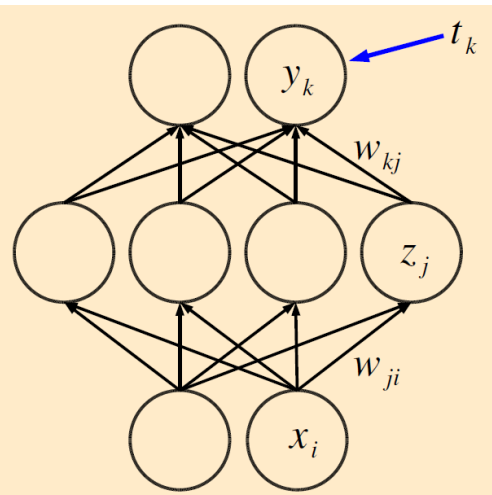
# Again, it's “complex” but not incomputable!

The chain rule makes it all possible...

$$\frac{\partial E}{\partial w_{ij}} = \frac{\partial}{\partial w_{ij}} \frac{1}{2} \sum_{k=1}^c (y_k - t_k)^2 = \sum_{k=1}^c (y_k - t_k) \frac{\partial y_k}{\partial w_{ij}}$$

Chaining back the **deltas** through **weighted sums** works throughout the network!

$$\Delta w_{ij} \propto -\frac{\partial E}{\partial w_{ij}}$$



For output units...

$$\delta_j = g'(net_j)(a_j - t_j)$$

For hidden units...

$$\delta_j = g'(net_j) \sum_k w_{jk} \delta_k$$

$$\Delta w_{ij} \propto -\delta_j a_i$$

Note that these equations apply for any weight connecting unit  $i$  to unit  $j$ .

# Generalized Delta Rule

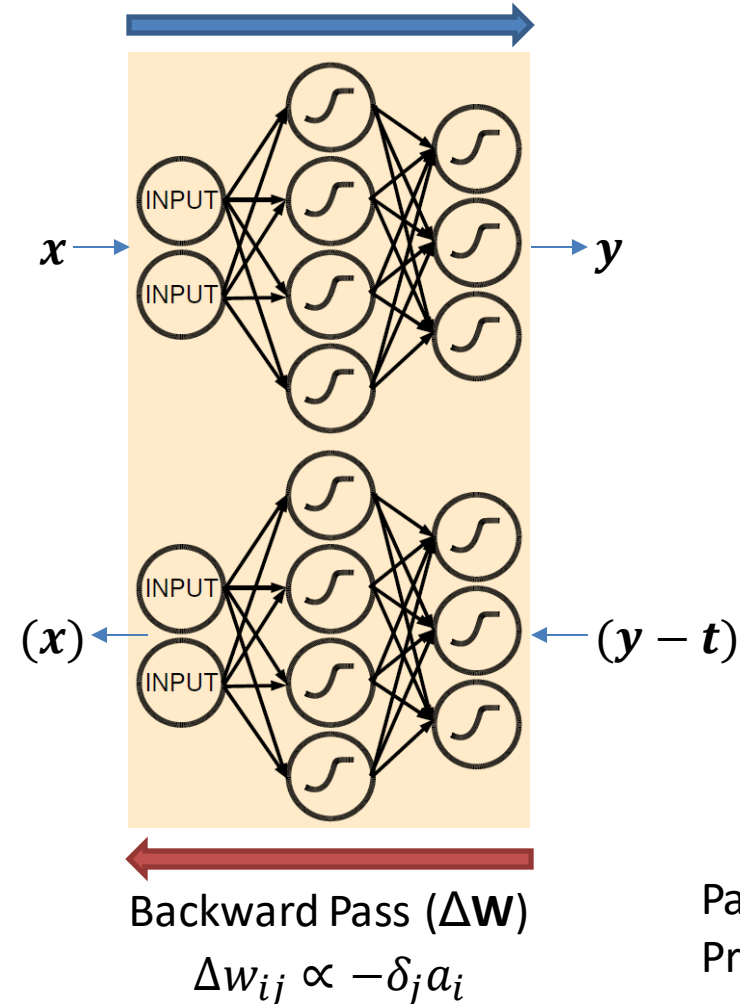
- Also known as **error backpropagation** or “**backprop**”
- Two step process:
  - **Prediction:** the inputs are provided and information flows to calculate the output activations (forward pass)
  - **Fitting:** errors are calculated at the output layer and information flows in reverse to calculate the weight updates (backward pass)
- Issue: biological plausibility
  - Real neurons do not pass information backward across synapses...

$$\text{Output unit: } \delta_j = g'(net_j)(a_j - t_j)$$

$$\text{Hidden unit: } \delta_j = g'(net_j) \sum_k w_{jk} \delta_k$$

$$a_i = g(net_i) \quad net_i = w_{i0} + \sum_j w_{ij} a_j$$

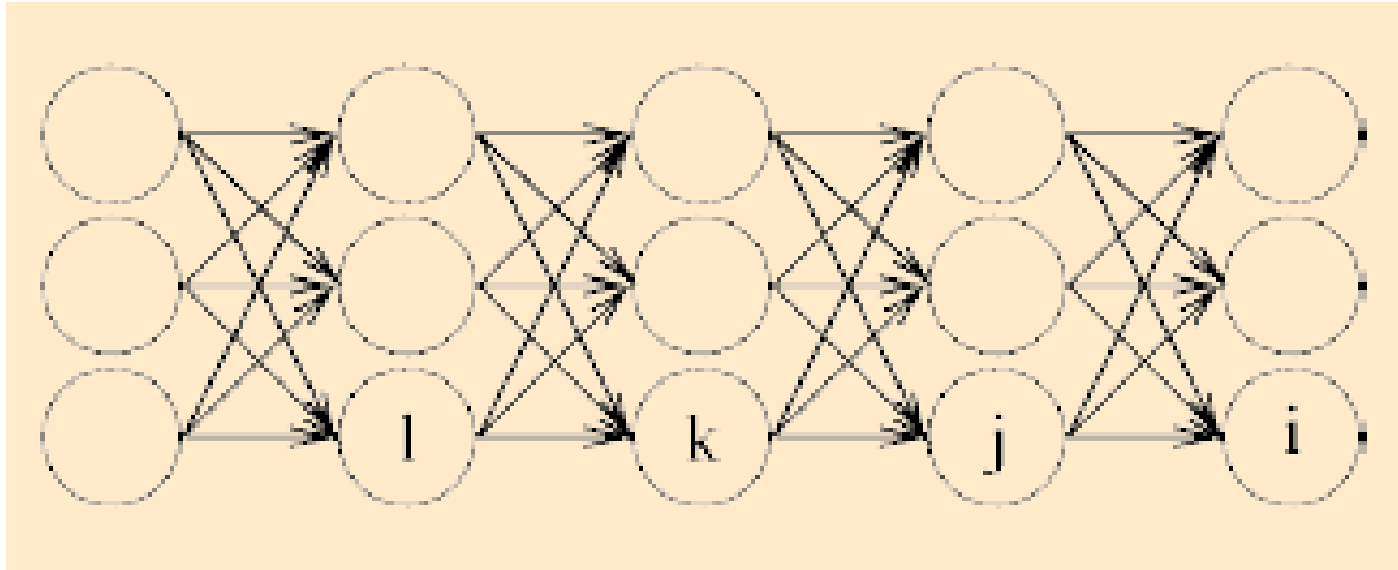
Forward Pass ( $\mathbf{W}$ )



Parallel Distributed  
Processing (Rumelhart  
and McClelland, 1986)

# Deep Learning?

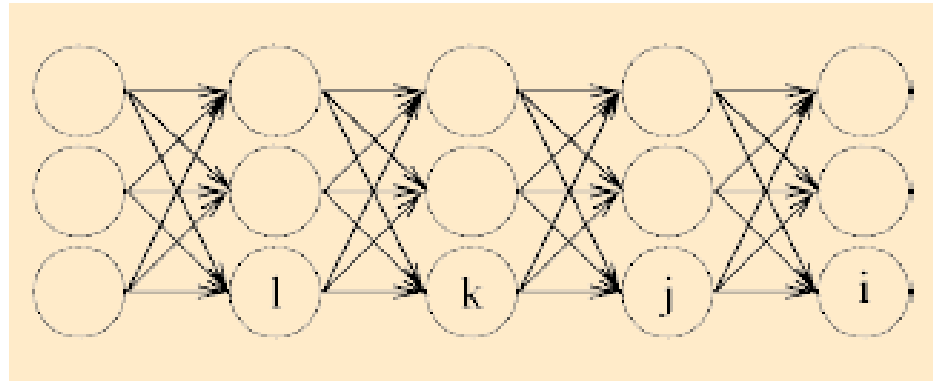
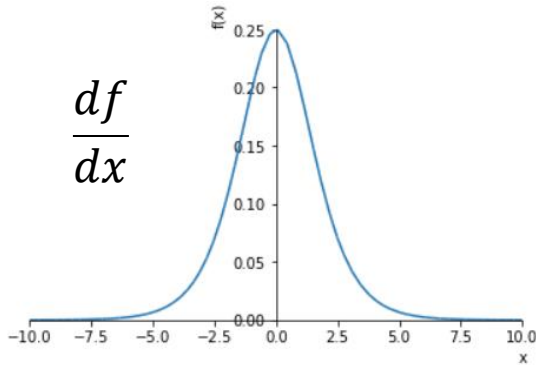
- Generalized Delta Rule or Error Backpropagation
  - This algorithm scales to an arbitrary number of hidden layers
  - Maybe...



$$\delta_j = g'(net_j) \sum_i w_{ij} \delta_i$$
$$\delta_k = g'(net_k) \sum_j w_{jk} \delta_j$$
$$\delta_l = g'(net_l) \sum_k w_{kl} \delta_k$$

# The Vanishing Gradient Problem

- Getting caught in the flat, planar regions of the error surface is problematic
- We **need** non-linear activation functions to make use of multiple layers, but typical non-linear activation functions (sigmoid, tanh) cause the gradient to **vanish**...

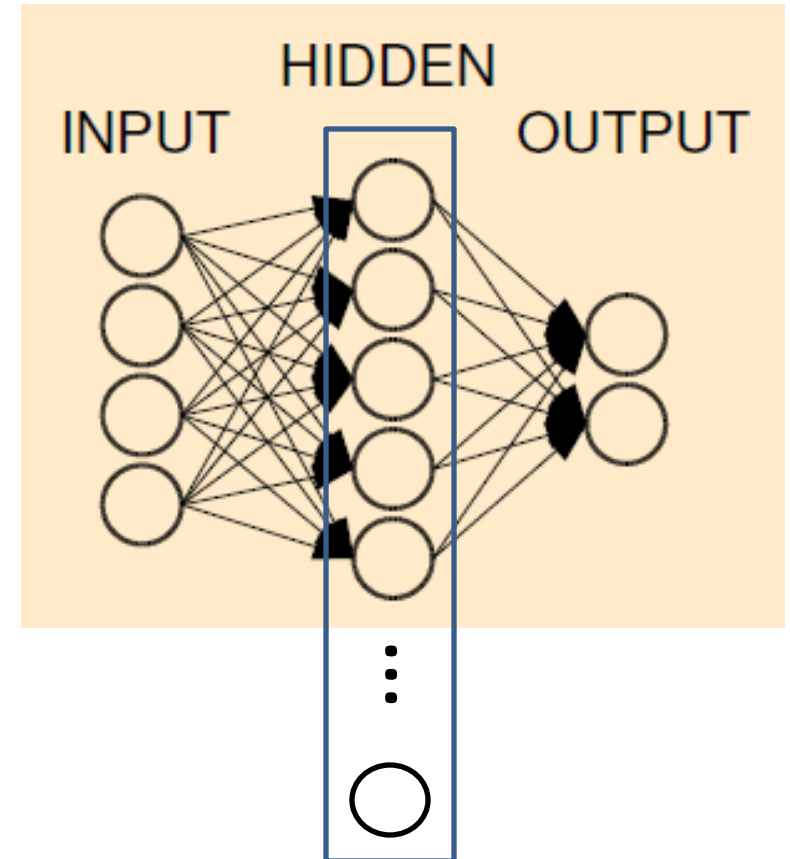


$$\delta_j = g'(net_j) \sum_i w_{ij} \delta_i$$
$$\delta_k = g'(net_k) \sum_j w_{jk} \delta_j$$
$$\delta_l = g'(net_l) \sum_k w_{kl} \delta_k$$

- We will see some workarounds to this *particular* issue soon, but even still, gradient optimization is tricky...

# So is wider better?

- Given the practical problems with the gradient disappearing the “wider is better” trend existed for over a decade without any clear resolution...
- Potential problems
  - (Major issue) Complex decision boundaries are probably better described by multilayered partitions instead of a single, very complex partition: might even be easier to learn this way if we could...
  - (Minor issue) Wider networks allow for poorer parallel training and performance since we can't process more than two batches of patterns through a feed-forward network at any one time



# Encodings Matter

- Let's reencode the XOR problem vectors...
  - Let each value (0 or 1) be replaced with a corresponding vector
    - zero=[0 1]
    - one=[1 0]
  - For each pattern:
    - compute the *outer product* of the value vector for column 1 and the value vector for column 2
    - Flatten into a 4 element vector
  - Combine vectors into new input pattern matrix...
- Even Hebbian learning can solve this!

$$\begin{bmatrix} 0 & 0 \\ 0 & 1 \\ 1 & 0 \\ 1 & 1 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$

First pattern is [0 0], so use two vectors: [0 1] and [0 1]

$$\begin{matrix} \times & 0 & 1 \\ 0 & \begin{bmatrix} 0 & 0 \end{bmatrix} \\ 1 & \begin{bmatrix} 0 & 1 \end{bmatrix} \end{matrix}$$

Flatten into a four element vector...

$$\begin{bmatrix} 0 & 0 & 0 & 1 \end{bmatrix}$$

Do the same for all four:

$$\begin{bmatrix} 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 0 & 1 & 0 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix} = \begin{bmatrix} 0 \\ 1 \\ 1 \\ 0 \end{bmatrix}$$



# Higher-dimensional Projections

---

- **Projecting data into higher dimensions** makes *decision boundaries* (or functional shapes for regressions) *easier to obtain...*
- Generally, the more units used, the more complex this projection *may* be
- However, this *kind* of solution is limited in the sense that it has to be performed in this specific way (no **nested, repeated, functional, or relational** solutions...)

